

AHMAD MUJTABA

GenAI Engineer | Agentic AI Specialist | Data Scientist

 +91-98070-52640  ahmad.iiitk@gmail.com  LinkedIn  GitHub

January 29, 2026

To the Hiring Manager,

In the current landscape of Artificial Intelligence, the challenge is no longer just accessing a Large Language Model—it is orchestrating it to behave reliably in an enterprise environment. As a **GenAI Engineer** and **Data Scientist** currently working with Deloitte US-India, I specialize in the "last mile" of AI engineering: moving systems from experimental demos to deterministic, production-grade workflows.

My core expertise lies in **Agentic AI**, **Model Context Protocol (MCP)**, and **Eval-Driven Development**. I don't just build chatbots; I build autonomous systems that execute complex business logic with measurable accuracy.

At Deloitte, I am currently tackling the reliability issues inherent in autonomous agents. When out-of-the-box **Computer-Using Agents (CUA)** failed to navigate complex business UIs (achieving only ~38% success), I did not simply prompt harder. Instead, I architected a **RAG-driven orchestration layer** using Milvus to inject granular Standard Operating Procedures (SOPs) dynamically into the agent's context. This engineering-first approach grounded the agent's reasoning, boosting task automation success rates to ~80% and significantly reducing human workload.

Furthermore, I understand that "generative" capabilities are often a liability in compliance-heavy sectors like Insurance. To solve this, I designed a **Reasoning-Driven IDP pipeline** for unstructured policy documents that achieved **99% extraction accuracy**. By implementing a unique "Canonical Reconciliation" loop—where the model self-corrects against authoritative reference docs—I proved that GenAI can be made safe enough for rigorous downstream analytics.

My technical foundation is not limited to modeling; it extends to full-stack deployment. Whether it is optimizing token usage by 40% using **MCP** and **Playwright**, building scalable ETL pipelines in **Azure Databricks (PySpark)**, or containerizing workflows with **Docker** on AWS, I ensure that the AI solutions I build are reproducible, scalable, and cost-effective.

I am eager to bring this rigorous, system-level approach to AI engineering to your team. I look forward to discussing how my experience in building self-healing agents and robust RAG architectures can drive value for your organization.

Sincerely,

Ahmad Mujtaba